

EXPERIMENT - 4

AIM

Using the Employee schema, write and execute SQL queries with INNER JOIN, LEFT JOIN, self-join, 3-way join, correlated subqueries, EXISTS, and simulated INTERSECT and EXCEPT. Compare the execution plans of selected queries using EXPLAIN.

OBJECTIVES

- To understand how different types of JOIN operations combine rows from related tables.
- To perform a self-join on the Employee table to compare employees within the same department.
- To use correlated subqueries and EXISTS for row-by-row and existence-based filtering.
- To simulate INTERSECT and EXCEPT operations using subqueries supported by MySQL.
- To inspect and compare query execution plans using EXPLAIN.

DATABASE USED

The queries use the CompanyDB database and the Employee–Department–Project schema from the reference experiment. The reference document defines 30 employees, 5 departments and 8 projects. The Employee table contains emp_id, emp_name, salary, hire_date, dept_id and project_id, with foreign keys to Department and Project.

USE CompanyDB;

SCHEMA REFERENCE

Table	Important Columns	Relationship
Employee	emp_id, emp_name, salary, hire_date, dept_id, project_id	dept_id → Department; project_id → Project
Department	dept_id, dept_name	Primary key: dept_id
Project	project_id, project_name, budget, dept_id	Primary key: project_id; dept_id → Department

The source document creates Department, Project and Employee in this order and establishes the foreign-key relationships shown above.

SOURCE CODE / SQL QUERIES

1. INNER JOIN

Display employee names, salaries and their department names using an INNER JOIN.

```
SELECT
    e.emp_id,
    e.emp_name,
    e.salary,
    d.dept_name
FROM Employee e
INNER JOIN Department d
    ON e.dept_id = d.dept_id
ORDER BY e.emp_id;
```

Output

 SHIVESH KUMAR KANOUIA

I/O AI Agent 87 ms 

STDIN

Input for the program (Optional)

Output:

emp_id	emp_name	salary	dept_name
1	Aarav	65000.00	IT
2	Vivaan	72000.00	IT
3	Aditya	58000.00	IT
4	Arjun	81000.00	IT
5	Kabir	69000.00	IT
6	Reyansh	62000.00	IT
7	Ananya	52000.00	HR
8	Diya	48000.00	HR
9	Myra	61000.00	HR
10	Sara	55000.00	HR
11	Ishita	47000.00	HR
12	Meera	59000.00	HR
13	Rohan	75000.00	Finance
14	Karan	68000.00	Finance
15	Nikhil	83000.00	Finance
16	Yash	62000.00	Finance
17	Manav	71000.00	Finance
18	Dev	56000.00	Finance
19	Aanya	54000.00	Marketing

The full query returns all 30 employees. The Department values come from the Department table in the reference data.

2. LEFT JOIN

Display all departments and the employees belonging to them. LEFT JOIN keeps every department from the left table even if no matching employee exists.

```
SELECT
  d.dept_id,
  d.dept_name,
  e.emp_name,
  e.salary
FROM Department d
LEFT JOIN Employee e
  ON d.dept_id = e.dept_id
ORDER BY d.dept_id, e.emp_id;
```

OUTPUT

 SHIVESH KUMAR KANOUIJIA

I/O

AI Agent

93 ms 

STDIN

Input for the program (Optional)

Output:

dept_id	dept_name	emp_name	salary
1	IT	Aarav	65000.00
1	IT	Vivaan	72000.00
1	IT	Aditya	58000.00
1	IT	Arjun	81000.00
1	IT	Kabir	69000.00
1	IT	Reyansh	62000.00
2	HR	Ananya	52000.00
2	HR	Diya	48000.00
2	HR	Myra	61000.00
2	HR	Sara	55000.00
2	HR	Ishita	47000.00
2	HR	Meera	59000.00
3	Finance	Rohan	75000.00
3	Finance	Karan	68000.00
3	Finance	Nikhil	83000.00
3	Finance	Yash	62000.00
3	Finance	Manav	71000.00
3	Finance	Dev	56000.00

In the supplied data, all five departments have employees, so no NULL employee row appears. The LEFT JOIN would show a department with NULL employee columns if such a department existed.

3. SELF-JOIN

Compare two employees belonging to the same department. The Employee table is joined to itself using two aliases.

```
SELECT
  e1.emp_name AS employee_1,
  e2.emp_name AS employee_2,
  e1.dept_id,
  e1.salary AS salary_1,
  e2.salary AS salary_2
FROM Employee e1
INNER JOIN Employee e2
  ON e1.dept_id = e2.dept_id
  AND e1.emp_id < e2.emp_id
ORDER BY e1.dept_id, e1.emp_id, e2.emp_id
LIMIT 10;
```

OUTPUT (first 10 rows)

 SHIVESH KUMAR KANOUIA

I/O

AI Agent

78 ms 

STDIN

Input for the program (Optional)

Output:

employee_1	employee_2	dept_id	salary_1	salary_2
Aarav	Vivaan	1	65000.00	72000.00
Aarav	Aditya	1	65000.00	58000.00
Aarav	Arjun	1	65000.00	81000.00
Aarav	Kabir	1	65000.00	69000.00
Aarav	Reyansh	1	65000.00	62000.00
Vivaan	Aditya	1	72000.00	58000.00
Vivaan	Arjun	1	72000.00	81000.00
Vivaan	Kabir	1	72000.00	69000.00
Vivaan	Reyansh	1	72000.00	62000.00
Aditya	Arjun	1	58000.00	81000.00

The condition `e1.emp_id < e2.emp_id` prevents the same pair from being repeated in reverse order and prevents an employee from being paired with itself.

4. 3-WAY JOIN

Display employee name, department, project and salary by joining Employee, Department and Project.

```
SELECT
    e.emp_name,
    d.dept_name,
    p.project_name,
    e.salary
FROM Employee e
INNER JOIN Department d
    ON e.dept_id = d.dept_id
INNER JOIN Project p
    ON e.project_id = p.project_id
ORDER BY d.dept_name, e.emp_name;
```

OUTPUT

 SHIVESH KUMAR KANOUIA

I/O

AI Agent

111 ms 

STDIN

Input for the program (Optional)

Output:

emp_name	dept_name	project_name	salary
Dev	Finance	Financial Dashboard	56000.00
Karan	Finance	Financial Dashboard	68000.00
Manav	Finance	Financial Dashboard	71000.00
Nikhil	Finance	Financial Dashboard	83000.00
Rohan	Finance	Financial Dashboard	75000.00
Yash	Finance	Financial Dashboard	62000.00
Ananya	HR	Recruitment Portal	52000.00
Diya	HR	Employee Wellness	48000.00
Ishita	HR	Recruitment Portal	47000.00
Meera	HR	Employee Wellness	59000.00
Myra	HR	Recruitment Portal	61000.00
Sara	HR	Employee Wellness	55000.00
Aarav	IT	Cloud Migration	65000.00
Aditya	IT	Mobile Application	58000.00
Arjun	IT	Cloud Migration	81000.00
Kabir	IT	AI Analytics	69000.00

The reference experiment uses the same three-table join pattern and orders the result by department and employee name.

5. CORRELATED SUBQUERY

Display employees whose salary is greater than the average salary of their own department. The inner query refers to the outer Employee row, making it a correlated subquery.

```
SELECT
  e.emp_id,
  e.emp_name,
  e.salary,
  e.dept_id
FROM Employee e
WHERE e.salary > (
  SELECT AVG(e2.salary)
  FROM Employee e2
  WHERE e2.dept_id = e.dept_id
)
ORDER BY e.dept_id, e.salary DESC;
```

OUTPUT

 SHIVESH KUMAR KANOUIA

I/O

AI Agent

112 ms 

STDIN

Input for the program (Optional)

Output:

emp_id	emp_name	salary	dept_id
4	Arjun	81000.00	1
2	Vivaan	72000.00	1
5	Kabir	69000.00	1
9	Myra	61000.00	2
12	Meera	59000.00	2
10	Sara	55000.00	2
15	Nikhil	83000.00	3
13	Rohan	75000.00	3
17	Manav	71000.00	3
24	Navya	70000.00	4
22	Riya	66000.00	4
20	Kiara	63000.00	4
29	Rudra	76000.00	5
26	Dhruv	73000.00	5
27	Atharv	67000.00	5

The department averages used by this query are consistent with the reference experiment's GROUP BY/HAVING results, including IT 67833.33, Finance 69166.67 and Operations 66500.00. □

6. EXISTS

Find departments that have at least one employee with a salary greater than 80,000.

```
SELECT
  d.dept_id,
  d.dept_name
FROM Department d
WHERE EXISTS (
  SELECT 1
  FROM Employee e
  WHERE e.dept_id = d.dept_id
  AND e.salary > 80000
)
ORDER BY d.dept_id;
```

OUTPUT

 SHIVESH KUMAR KANOUIA

 I/O  AI Agent 89 ms 

STDIN

Input for the program (Optional) 

Output:

dept_id	dept_name
1	IT
3	Finance

EXISTS returns TRUE as soon as the subquery finds at least one qualifying employee for the current department.

7. SIMULATED INTERSECT

MySQL can express the intersection of two employee sets by requiring the employee ID to occur in both subqueries. Here, the two sets are: employees in IT and employees whose salary is greater than 70,000.

```
SELECT emp_id, emp_name, salary
FROM Employee
WHERE emp_id IN (
  SELECT emp_id
  FROM Employee
  WHERE dept_id = 1
)
AND emp_id IN (
  SELECT emp_id
```

```

FROM Employee
WHERE salary > 70000
)
ORDER BY emp_id;

```




SHIVESH KUMAR KANOUIA

I/O
AI Agent
92 ms 

STDIN

Input for the program (Optional)

Output:

```

+-----+-----+-----+
| emp_id | emp_name | salary |
+-----+-----+-----+
|      2 | Vivaan   | 72000.00 |
|      4 | Arjun    | 81000.00 |
+-----+-----+-----+

```

The result contains only employees that satisfy both sets, which is the behavior of an INTERSECT operation.

8. SIMULATED EXCEPT

Simulate EXCEPT by selecting employees with salary greater than 70,000 while excluding employees belonging to IT.

```

SELECT emp_id, emp_name, salary
FROM Employee
WHERE emp_id IN (
    SELECT emp_id
    FROM Employee
    WHERE salary > 70000
)
AND emp_id NOT IN (
    SELECT emp_id
    FROM Employee
    WHERE dept_id = 1
) ORDER BY emp_id;

```




SHIVESH KUMAR KANOUIA

I/O
AI Agent
87 ms 

STDIN

Input for the program (Optional)

Output:

```

+-----+-----+-----+
| emp_id | emp_name | salary |
+-----+-----+-----+
|      13 | Rohan    | 75000.00 |
|      15 | Nikhil   | 83000.00 |
|      17 | Manav    | 71000.00 |
|      26 | Dhruv    | 73000.00 |
|      29 | Rudra    | 76000.00 |
+-----+-----+-----+

```

9. EXPLAIN – EXECUTION PLAN

EXPLAIN shows how the MySQL optimizer intends to execute a SELECT statement. It can be used to inspect access type, possible indexes, selected index, estimated rows and additional operations.

A. EXPLAIN for INNER JOIN

```
EXPLAIN
SELECT
    e.emp_name,
    d.dept_name
FROM Employee e
INNER JOIN Department d
    ON e.dept_id = d.dept_id
WHERE e.salary > 70000;
```

B. EXPLAIN for Correlated Subquery

```
EXPLAIN
SELECT
    e.emp_name,
    e.salary
FROM Employee e
WHERE e.salary > (
    SELECT AVG(e2.salary)
    FROM Employee e2
    WHERE e2.dept_id = e.dept_id
);
```

C. EXPLAIN for EXISTS

```
EXPLAIN
SELECT d.dept_name
FROM Department d
WHERE EXISTS (
    SELECT 1
    FROM Employee e
    WHERE e.dept_id = d.dept_id
    AND e.salary > 80000
);
```

10. COMPARISON OF EXECUTION PLANS

Query	Main operation	Plan aspect to compare	Expected observation
INNER JOIN	Employee ↔ Department	Join access on dept_id	Checks how Employee and Department are accessed and joined.
Correlated subquery	Outer Employee + repeated department AVG logic	Subquery execution and rows examined	May require more work because the inner query is logically related to each outer row.
EXISTS	Department + existence test on Employee	Early existence check and filtering	Can stop searching once a qualifying Employee row is found.

For a small dataset such as the 30 employees in this experiment, the optimizer may choose table scans because the cost is low. With larger tables and suitable indexes on join/filter columns, the execution plan can change. The comparison should therefore be based on the actual EXPLAIN output produced in the lab environment.

RESULT

The Employee–Department–Project database was used to successfully demonstrate INNER JOIN, LEFT JOIN, self-join, 3-way JOIN, correlated subqueries, EXISTS, simulated INTERSECT and simulated EXCEPT. EXPLAIN statements were also prepared to compare how MySQL plans the different query forms.

CONCLUSION

Different SQL constructs can produce related results while using different execution strategies. JOIN operations combine rows from tables, correlated subqueries evaluate related values for outer rows, EXISTS tests whether matching rows exist, and subqueries can simulate set operations such as INTERSECT and EXCEPT. EXPLAIN provides information needed to study and compare these execution strategies.